OMO RESEARCH · WHITEPAPER · OCTOBER 2026

Worlds, Not Windows

Spatially aware creative agentic interfaces — findings so far

Harrison Edwards · Omo Research

Findings so far, from using it every day, live on stream, 25–29 October 2026. Not a controlled experiment.



REAL · LIVE STREAM, 30 OCT 2026 A hall in the underground base where most of the agent work lived, each place with its own screen. Cropped still from the author's stream, in the live world with real agent sessions.

01 The problem

AI agents no longer answer one question and stop. They run for hours. They write code, cut video, draft posts and wait for approval. And we run many of them at once.

Most of us still manage them through chat windows and terminal tabs. That works for one agent. It mostly works for three. With twenty, the tab bar stops telling you the things you need to know. Which session is which? Which one is busy? Which one finished an hour ago? Which one is stuck, waiting for you?

A tab is a label in a row of labels. It looks the same whether the agent behind it is working hard, done, or blocked on a question. You find out by clicking through them, one by one.

On the first night of streaming with Omo Space, Harrison put the alternative like this:

> "So it's like I'm spatially aware of my workspace. … It's not a chat window, it's an interface."
>
> [T1, 00:38:33–00:38:50, live]

02 The idea: worlds, not windows

Our idea is simple. **Give agent work a place.** Not a tab, not a window, but somewhere you can walk to.

People are good at remembering where things are. We remember which room the books are in and which shelf holds the tools. We are much worse at remembering that the video agent was the seventh tab from the left. So instead of a row of tabs, we put each agent session in a place in a world, with a screen you can read and a lamp that shows when it is working.

Places come at every scale. A single room for one job. A cave of rooms for a whole content pipeline. A tower, an office, an arena. A whole world for one line of work. The shape is chosen to fit the work, not imposed by the software.

And the places do not have to be built by hand. In Omo Space, coding agents build most of them from short requests. You ask for a place and an agent builds it. That changes the cost of a spatial interface: making a new room is closer to opening a new folder than to building a level in a game.



ILLUSTRATION The idea at its largest scale: an organisation drawn as an open world of districts, with agents at work in each. A generated illustration, not a screenshot. It shows where we would like to go, not what we have seen.

03 How Omo Space works

Omo Space is a Fabric mod for Minecraft Java 26.2, plus a small local bridge. It puts real agent sessions (Hermes Agent and Claude Code) and live web pages on screens inside the game.

The key design choice: **the game does not run the agents.** Each workspace is a named terminal session (tmux) on the user's own computer. It has one or more tabs, each running an agent or a plain shell. A bridge that listens only on the local machine paints those sessions onto in-world screens and sends your keys back. Minecraft is the view and the controls. The work lives outside it.

IN THE WORLD	WHAT IT IS ON THE COMPUTER
Hermes Computer	A named workspace whose first tab runs a Hermes agent in its own profile (its own memory, skills and model)
Terminal	A workspace tab running a shell or another agent CLI
Monitor wall / Slim Screen	A larger view of a tab, or a pinned live web page
Linker	Binds a screen to a workspace from a distance; several screens can mirror one
Status Lamp	Gold if any tab in the workspace is busy, green when idle, red once interrupted, dark with no live workspace
Name Plate	The workspace name, readable from across the room

A place is anything you build around these objects: a room, a cave, a tower, a world. Places are just builds. They carry no permissions and run nothing.

What survives what

Because the agents live in tmux, not in the game, closing things in the game does not stop them. Each Hermes tab is one named conversation, so reopening a screen brings back the same conversation, not a new one.

WHAT HAPPENS	RUNNING JOBS	HERMES CONVERSATION
Close the screen, walk away	Keep running	Keeps running
Quit or crash Minecraft	Keep running	Keeps running
Restart the bridge	Keep running	Keeps running
Hermes tab idle and unseen 20 min	—	Naps; wakes when shown again
Restart the PC	Stopped	Saved; tab shows INTERRUPTED; press Resume

So a PC restart does stop running work. Saved conversations come back with a button. Nothing re-runs on its own: in the documentation's words, "Nothing is ever re-run without you" [1]. This table is the documented behaviour. Web pages pinned on walls are a separate kind of state.

Two things the system does not do yet. A lamp shows activity, but there is no "waiting for you" state. And agents do not pass work to each other through the world. Any pipeline that runs from place to place is the user's own scripts and agent

instructions.



REAL · LIVE STREAM, 30 OCT 2026 A Hermes session on a wall terminal (right), with other agents' screens across the room. Cropped still; face camera, stream label and game HUD removed.

04 What we've seen so far

Everything here comes from Harrison using Omo Space as his main workspace for five days, live on stream, 25–29 October 2026: about 29 hours of transcribed streams. He ran a real content pipeline (ideas, references, recording, editing, publishing, analytics) with roughly one agent per stage. A friend also built his own world with it and showed it on stream.

We saw place used in four ways: to **find**, to **resume**, to **monitor** and to **orchestrate**. The quotes are what was said on stream. They are reports, not measurements.

Find: walk to it

Every workspace sits in a fixed, named place. You remember where an agent is, not which tab it was in. Live on the first night, Harrison named the problem the world was built to fix:

> “Getting lost with all these tabs and things” [T1, 01:13:45, live]

Later he reported about 25 sessions running at once, “within the interface that we can conceptually understand” [T6, 07:37:22, live].



REAL · LIVE STREAM, 30 OCT 2026 Find: a corridor of places, each with its own screen. The image shows layout, not an outcome.

Resume: come back to it

The agents run outside the game, so leaving the game does not stop them, and the place is still there when you return. After quitting and reloading on the first night:

> “So we can close it enjoy our day And then Come back Boom the agents are still working right here” [T1, 02:12:55, live]

On 29 October he said agents had kept editing, about 22 clips, after he closed Minecraft [T8, 00:06:30–00:08:38, live]. That fits the table above: closing the game does not stop jobs. It does not mean work ran while the PC was off; a restart stops running jobs. He also tied this to how it felt to start work: “I’m not dreading any pain / Or like context loss because I know anytime I want to work on the robotics project / I just fly down here” [T4, 02:21:42, live].

Monitor: glance at lit screens

Lamps and lit screens let you read activity from across a room. The friend described exactly this cue, live, in his own world:

> “How I can tell when the task is like running because their screen will turn on”
>
> [T8, 00:48:43, live; a friend, speaker inferred]

He also said watching the team was more fun than “just looking at you know Claude you know just typing” [T8, 01:00:57, live; a friend, speaker inferred].

Lamps show who is busy. They do not yet show that an agent is waiting for an answer; that is the next signal to build.



REAL · LIVE STREAM, 30 OCT 2026 Monitor: lit screens readable from the cavern entrance. The image shows layout, not an outcome.

Orchestrate: move work room to room

Harrison laid his content pipeline out as a row of places, about one stage and one agent each:

> “We just go to that room and just start building / And when we’re done, then we just leave that room” [T2, 01:04:14, live]

By 29 October he wanted a place for each to-do item, “a repeatable pattern” [T8, 03:28:22, live].

The hand-offs between rooms are his own scripts and agent instructions. Laying the pipeline out in space made its stages visible and easy to visit.

Places at every scale

The places were not just rooms. The pipeline lived in a cave of five rooms. He wanted whole worlds for lines of work: “We can build a whole world for the paid ads” [T7, 02:48:38, live]. To show how personal it could get, he imagined a pirate ship holding your tax work; that one stayed an idea. He said agents did most of the building: “I didn’t build any of this” [T6, 07:45:10, live].



REAL · LIVE STREAM, 30 OCT 2026 An atrium.



REAL · LIVE STREAM, 30 OCT 2026 A hall, with an embodied agent.

A friend took the idea his own way. New to Hermes and on Windows, he described about 24 hours of work: “It just gave Claude the GitHub and then gave it a few prompts and then it just it built the world” [T8, 01:04:44, live; a friend, speaker inferred]. His world had an AI office with a manager agent, teams, a podium and a sports arena; those were his own additions, not part of Omo Space.

05 Web pages with a terminal attached

Late in the week Harrison described the core pattern in its most general form, apart from Minecraft:

> “the concept of spatial interfaces with persistent webpages is easy to use and talk to and control. … We’re doing it inside Minecraft, but I believe there is not enough … attention, these persistent webpages with a terminal attached to it. And with this, you can build whatever pipeline that you want.” [T8, 03:38:57–03:39:18, live]

Pieced together from several moments, the pattern goes like this. You ask an agent at a station for a page for one job: a review queue, an analytics view, a publishing form. You pin it on a wall in the right place. You check it, and you approve or ask for changes through the terminal next to it. Then you come back to it later.

The page does the job-specific work. The place gives it a fixed home next to related work. The terminal beside it means the page is never finished; you can always ask the agent to change it. An ordinary browser tab could host the same page, and we have not separated those two effects. The in-world browser has limits too: drop-down menus could not be selected [T8, 03:22:45, live].



REAL · LIVE STREAM, 30 OCT 2026 A page with a terminal attached: the review stage of the video pipeline, with approve, send back and move-to-done controls, beside the editing agent’s terminal.

06 Design lessons

- **Give work a place.** A persistent, recognisable place per line of work, with room for several sessions per place, mirrored screens and any scale.
- **Separate address from state.** A place is an address. The process, the conversation, the page and the world behind it each need their own stated guarantee, and each needs testing.
- **Activity is not "waiting for you".** A lit screen says busy. A question needs its own state, and a signal that follows the person into other rooms. Do not rely on colour alone.
- **Make hand-offs explicit.** Moving work between places should be a visible object or action, not something implied by layout.
- **Let the work choose the shape.** Rooms, caves, towers, ships, whole worlds; built by agents or by hand.
- **Honest state.** Never show progress that is not happening. Re-run nothing on the user's behalf. Make the cost of the world itself visible, not only the work it saves.

07 What's next

We do not claim worlds beat tabs. We have not measured that. Here is what we'd like to measure next.

- **Finding.** How fast can someone find a named agent among 12 or 24 sessions in a world, compared with tabs? The tab side should be fair: same agents, same persistence, good names, status badges and search, so any difference comes from space and not from better tools.
- **Resuming.** After a day away, how long until the first useful action, in a world versus tabs?
- **A "waiting for you" signal.** We have not built it yet. Once it exists: is a blocked agent noticed sooner? And does it still work when the agent is in another room, or only with a signal that follows you?
- **One-prompt places.** Logging every place an agent builds, including failed and abandoned ones, to see how many get used as-is, how many need fixing, and how many are still in use two weeks later.

Some of the raw material already exists: bridge logs of sleep, resume and interruption events, the Hermes session stores, and the recordings.

08 Prior work

Using space to organise computer work is an old idea. Rooms (1986) split a crowded screen into virtual rooms [2]. Data Mountain (1998) showed people can use spatial memory to place and find documents [3]. The Task Gallery (2000) hung running applications on the walls of a 3D gallery [4]. Calm technology argued for displays that show state "without being obtrusive" [5]. All of that was about documents, windows and tasks, not agents that keep working on their own. The closest recent work tackles agent supervision with dashboards: ParallelPilot adds "an ambient dashboard alongside existing coding tools" [6], and AgentGUI is an interface for observing and steering long-running agents [7]. Both report controlled user experiments; we have none. What Omo Space adds is a combination: places, built by people or by agents, that hold long-running sessions and agent-made pages.

References

1. Omo Research. 2026. Omo Space documentation: Agents and terminals. docs/AGENTS-AND-TERMINALS.md, revision 75d4677. https://github.com/omoresearch/omo-space
2. D. A. Henderson Jr. and S. K. Card. 1986. Rooms: The use of multiple virtual workspaces to reduce space contention in a window-based graphical user interface. *ACM Transactions on Graphics*. https://doi.org/10.1145/24054.24056
3. G. Robertson, M. Czerwinski, K. Larson, D. Robbins, D. Thiel and M. van Dantzich. 1998. Data Mountain: Using spatial memory for document management. In *Proc. UIST '98*. ACM. https://www.microsoft.com/en-us/research/publication/data-mountain-using-spatial-memory-for-document-management
4. G. Robertson et al. 2000. The Task Gallery: A 3D window manager. In *Proc. CHI '00*. ACM. https://doi.org/10.1145/332040.332482
5. M. Weiser and J. S. Brown. 1995. Designing calm technology. Xerox PARC, 21 December 1995. https://calmtech.com/papers/designing-calm-technology
6. T. Long, W. Shi, H. Mozannar, M. Murad and R. Hosn. 2026. ParallelPilot: Supporting coordination and monitoring in parallel AI coding. arXiv:2609.33113. https://arxiv.org/abs/2609.33113
7. X. Zhao, J. Sohn, Q. Zheng and M. Moor. 2026. AgentGUI: An interface for observing and steering long-running AI agents. arXiv:2607.26300. https://arxiv.org/abs/2607.26300

– About this paper

These are findings from the author's own use of a system he built, not a controlled study. One person used it as his main workspace, live on stream, and a friend built his own world with it. There is no comparison with tabs and nothing was measured. The longer research paper, *Worlds, Not Windows* (preprint draft v1.7), has the full method, episodes and limitations.

Quotes come from the stream transcripts, which were made automatically (Whisper) and may contain errors. Transcripts have no speaker labels; where a quote is from the friend, the speaker is inferred from

context. Locators such as [T1, 02:12:55] name the transcript segment and time; every quoted span was checked by script against that transcript. Live means said while working.

Images are labelled by source. Real · live stream images are cropped stills from the author’s stream of 30 October 2026, the day after the period described, with real agent sessions; they show what the places look like, not outcomes. The one Illustration is generated and is not a screenshot.